

CPSC 532P / LING 530A: Deep Learning for Natural Language Processing (DL-NLP)

Muhammad Abdul-Mageed

muhammad.mageed@ubc.ca

Natural Language Processing Lab

The University of British Columbia

Table of Contents

- 1 Problems With Gradients
- 2 Gated Recurrent Units (GRU)
- 3 Long-Short Term Memory Networks (LSTMs)

Vanishing and Exploding Gradients




Figure: A very long sequence, modelled with an RNN. Gradients can vanish and we will not know which gender a pronoun should be (male or female, e.g., in an MT task), or to which entity the pronoun refers (boy or girl)

Gradient Problems

- Gradients can **explode**, in which case we can clip them.
- Gradients can also **vanish**, which is a more serious problem.

Solving Long-Term Dependencies

Solutions For Gradient Problems

- Long-Short Term Memory (**LSTM**) networks introduced to solve the problem of long-term dependencies
- Gated Recurrent Units (**GRU**) (Cho et al., 2014; Chung et al., 2014): Simplification of LSTMs.
- We will **introduce GRUs first**, as they are simpler
- Some notation for LSTM modified from Andrew Ng, for pedagogical simplicity

Introducing a Memory Cell

augment network with a memory cell **C**




Figure: We will augment the network with a memory cell

Gradient Problems

- The **memory cell** will help us retain information over long sequences
- For example, we can still know **we need pronoun she**, maintaining the **female gender** (and retaining the correct **reference** to "the girl")

1: Simple GRU Cell

$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W_x \cdot [h_{t-1}, x_t])$$

$$h_t = z_t * \tilde{h}_t + (1 - z_t) * h_{t-1}$$

GRU Cells Notation Translation Table

- z : Update gate
- \tilde{h} : New candidate memory
- h_t : GRU output

2: GRU Cell

$$z_t = \sigma(W_z.[h_{t-1}, x_t])$$

$$r_t = \sigma(W_r.[h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W_x.[r_t * h_{t-1}, x_t])$$

$$h_t = z_t * \tilde{h}_t + (1 - z_t) * h_{t-1}$$

GRU Cells Notation Translation Table

- **r**: Relevance state (*reset gate*)

3: GRU With Bias

$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t] + b_z)$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t] + b_r)$$

$$\tilde{h}_t = \tanh(W_x \cdot [r_t * h_{t-1}, x_t] + b_h)$$

$$h_t = z_t * \tilde{h}_t + (1 - z_t) * h_{t-1}$$

GRU Cells Notation Translation Table

- **z**: Update gate
- **r**: Relevance state (*reset gate*)
- \tilde{h} : New candidate memory
- h_t : GRU activation (output)

4: Simple GRU Update

$$h_t = z_t * \tilde{h}_t + (1 - z_t) * h_{t-1}$$

GRU Gates

- **z**: **Update gate** is result of a **sigmoid** (between 0 and 1)
- **z close to zero**: We multiply by \sim zero, so we update candidate \tilde{h} very little (almost keep value of old memory cell h_t)
- **z close to one**: We multiply by \sim 1 and subtract by \sim 1, so old cell becomes almost equal to candidate)

How is LSTM Different?

5: Recall: GRU

$$z_t = \sigma(W_z.[h_{t-1}, x_t] + b_z)$$

$$r_t = \sigma(W_r.[h_{t-1}, x_t] + b_r)$$

$$\tilde{h}_t = \tanh(W_h.[r_t * h_{t-1}, x_t] + b_h)$$

$$h_t = z_t * \tilde{h}_t + (1 - z_t) * h_{t-1}$$

$$a_t = h_t$$

Note

- Cell activation a_t is the same as h_t .
- At each step, we start with $h_t = a_t$. They are different in LSTM.

Toward an LSTM

Changes

- Parts in **red** will change!
- $a_t \neq h_t$ and so we will use a_t
- We will also add new parts: forget gate and output gate ...

6: Recall: GRU

$$z_t = \sigma(W_z.[h_{t-1}, x_t] + b_z)$$

$$r_t = \sigma(W_r.[h_{t-1}, x_t] + b_r)$$

$$\tilde{h}_t = \tanh(W_x.[r_t * h_{t-1}, x_t] + b_h)$$

$$h_t = z_t * \tilde{h}_t + (1 - z_t) * h_{t-1}$$

$$a_t = h_t$$

Alternative Notation (Simpler)

GRU Cells Notation Translation Table

- **z**: Update state $\rightarrow \Gamma_u$
- **r**: Relevance state (*reset gate*) $\rightarrow \Gamma_r$
- \tilde{h} : New candidate memory cell state $\rightarrow \tilde{C}_t$
- h_t : GRU output $\rightarrow C_t$

LSTM: New Candidate Cell \tilde{h}_t

Updating \tilde{h}_t

- To acquire \tilde{h}_t , instead of h_{t-1} , we use the new a_{t-1} (since a_{t-1} is acquired differently than h_{t-1} , as we explain later)

7: New Candidate \tilde{C}_t

$$\tilde{h}_t = \tanh(W_x \cdot [a_{t-1}, x_t])$$

LSTM: New Candidate Cell \tilde{C}_t (New Notation)

Updating \tilde{C}_t

- To acquire \tilde{C}_t , instead of C_{t-1} , we use the new a_{t-1} (since a_{t-1} is acquired differently than C_{t-1} , as we explain later)

8: New Candidate \tilde{C}_t

$$\tilde{C}_t = \tanh(W_c.[a_{t-1}, x_t])$$

LSTM: Update and Forget Gates

Changes

- We will **not** use an **relevance gate** (Γ_r)
- Instead of using one update gate Γ_u , we will use **two gates** to control cell content: Γ_u (**update gate**), sometimes called **input gate** f_i , and Γ_f (**forget gate**)
- **Forget gate** will give the new memory cell C_t the option to keep or forget the **old cell** (C_{t-1}), but just add to it via **update gate** (Γ_u)

9: LSTM

$$\Gamma_u = \sigma(W_u \cdot [a_{t-1}, x_t])$$

$$\Gamma_f = \sigma(W_f \cdot [a_{t-1}, x_t])$$

$$C_t = \Gamma_u * \tilde{C}_t + \Gamma_f * C_{t-1}$$

Output Gate

- As mentioned, we will use an **output gate** (Γ_o)
- **Output gate** will enable us to update our a_t via element-wise multiplication by Γ_o

10: Output Gate

$$\Gamma_o = \sigma(W_o \cdot [a_{t-1}, x_t])$$

$$a_t = \Gamma_o * \tanh(C_t)$$

11: LSTM

$$\tilde{C}_t = \tanh(W_c \cdot [a_{t-1}, x_t])$$

$$\Gamma_u = \sigma(W_u \cdot [a_{t-1}, x_t])$$

$$\Gamma_f = \sigma(W_f \cdot [a_{t-1}, x_t])$$

$$\Gamma_o = \sigma(W_o \cdot [a_{t-1}, x_t])$$

$$C_t = \Gamma_u * \tilde{C}_t + \Gamma_f * C_{t-1}$$

$$a_t = \Gamma_o * \tanh(C_t)$$

12: LSTM

$$\tilde{C}_t = \tanh(W_c \cdot [a_{t-1}, x_t] + b_c)$$

$$\Gamma_u = \sigma(W_u \cdot [a_{t-1}, x_t] + b_u)$$

$$\Gamma_f = \sigma(W_f \cdot [a_{t-1}, x_t] + b_f)$$

$$\Gamma_o = \sigma(W_o \cdot [a_{t-1}, x_t] + b_o)$$

$$C_t = \Gamma_u * \tilde{C}_t + \Gamma_f * C_{t-1}$$

$$a_t = \Gamma_o * \tanh(C_t)$$

LSTM Schematic Illustrated

$$\tilde{C}_t = \tanh(W_c \cdot [a_{t-1}, x_t])$$

$$\Gamma_u = \sigma(W_u \cdot [a_{t-1}, x_t])$$

$$\Gamma_f = \sigma(W_f \cdot [a_{t-1}, x_t])$$

$$\Gamma_o = \sigma(W_o \cdot [a_{t-1}, x_t])$$

$$C_t = \Gamma_u * \tilde{C}_t + \Gamma_f * C_{t-1}$$

$$a_t = \Gamma_o * \tanh(C_t)$$

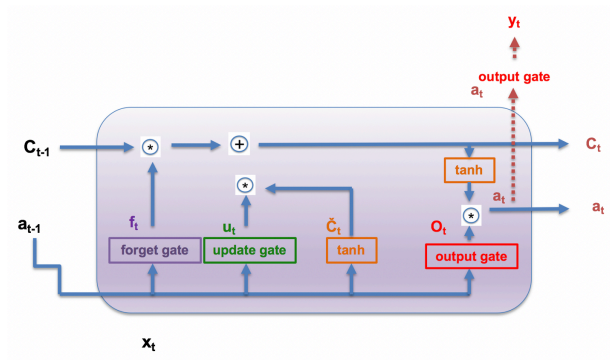


Figure: LSTM cell. [Inspired by Chris Olah]

Stacking LSTM Cells






Figure: LSTM cell.s stacked. Note: Each cell will need new-indexing (not shown in the Figure)